

## CMSC 422: Assignment 5 – Reinforcement Learning – Spring 2019

The goal of this project is to give you some experience using basic Q learning of sequential actions. You will be provided with the file *sarsaBook.py* – a basic implementation of the Sarsa Algorithm (from p. 243 of Marsland’s *ML* text) for learning how to traverse a simple “map” from location A to location F (Fig. 11.3, p. 234). It uses an  $\epsilon$ -greedy policy for action selection.

After familiarizing yourself with *sarsaBook.py*, copy it into a new file *qlearn.py*, renaming the function *SARSA()* to be *qlearn()*. Use *qlearn.py* for the rest of this problem. Begin by modifying *qlearn.py* in the following three ways:

Modify function *qlearn()* so that it implements the basic Q-Learning Algorithm (p. 242 of Marsland’s *ML* text) rather than the Sarsa Algorithm. Basically, this involves modifying the line starting with  $Q[s,a] += \mu * \dots$  near the end of the file.

Noting that the greedy action selection steps used in the original *sarsaBook()* function, for example  $a = \text{np.argmax}(Q[s,:])$ , can select illegal actions at times, alter the initialization of Q values in *qlearn()* so that this will never happen (i.e., so that Q is consistent with the network structure and does not allow selection of illegal actions by a pure greedy policy, but is otherwise initialized the same way).

Have *qlearn()* output the following information to a file *resultsQ.txt* during a run:

- immediately after initializing Q, write out the arrays R, t and Q (neatly formatted, 1 row per line)
- at the very end of a run, write out the final Q values (neatly formatted, one row per line)

Run the modified *qlearn()* function using the default parameters mu, gamma and epsilon that are embedded in the code, and then answer these questions based on the final Q values:

- Draw a picture of locations (nodes) and all legal actions (links), similar to that of Fig. 11.3 in Marsland’s *ML* text except now with directed links, where all of the learned Q values are written next to the corresponding links. For actions involving remaining at a location, write the Q values adjacent to the corresponding nodes.
- If an agent uses the learned Q values shown in (a) and a purely greedy action selection process, what would its sequence of actions be when starting an episode in state A?
- Was an optimal policy  $\pi$  learned? If so, why did the algorithm learn to follow the shortest path from A to the goal F, given that there are no explicit rewards built in for path shortness? If it did not learn an optimal policy, why not?
- Are the learned Q values good estimates of the discounted cumulative reward that would actually be received by an agent starting in an arbitrary city and following a purely greedy policy until task completion? Justify your answer.
- Given that the goal state F is supposed to be an absorbing state, what is wrong with the given transition matrix  $t$ ? (Don’t change  $t$ , just explain what is inconsistent with being an absorbing state.) Why does that not cause any significant problem at run time?

Copy *qlearn()* into a new file named *qdecay()*. Make two changes in *qdecay()* as follows. First, add the statement  $Q = 0.95 * Q$  at the beginning of each step of learning (right after the line ‘while s!=5:’ at the start of the inner while-loop), allowing the Q values to decay somewhat during each step of learning.

Second, have *qdecay()* output the same information as *qlearn()* except now to a file *resultsDecay.txt* . Run *qdecay()* using the same parameter values.

f. Draw a picture of locations (nodes) and all legal actions (links), similar to that in (a), where all of the learned Q values obtained using decay are written next to the corresponding links and nodes.

g. How does Q value decay like this affect the final Q values relative to those obtained without decay?

**What should I turn in?**

*The hardcopy and electronic submissions are due at different times.* The hard copy portion is due at the start of class Tuesday April 30, the electronic submission is due earlier at 11:30 pm Monday April 29.

*Hardcopy:* Your answers to the questions (a)-(g).

*Electronic submission:* Turn in a single zip file that includes your code in *qlearn()* and *qdecay()*, the result files *resultsQ.txt* and *resultsDecay.txt*, and any other files you created to run the code (e.g., a run script). Use the Computer Science Department project submission server that is located at <https://submit.cs.umd.edu> to submit this zip file.